



LUND UNIVERSITY

# **FMAN-45: Machine Learning**


## **Lecture 5: Support Vector Machines**

***Cristian Sminchisescu***

# Back to Binary Classification Setup...

- We are given a finite, possibly noisy, set of training data:  $\{\mathbf{x}_i, y_i\}, i = 1, \dots, N$ . Each input  $\mathbf{x}$  is paired with a binary output  $y$  (+1 or -1)
- Based only on training data, construct a machine that generates outputs  $y$ , given inputs  $\mathbf{x}$
- Now, a new sample is drawn from the same distribution as the training sample
- We wish to run the machine on the new sample input, and be able to classify it correctly, as either positive or negative


# Example: Face Detection



# Discriminant Function

Once again, we will restrict our attention to learning machines that separate the positive and negative examples using a linear function, with parameters  $(\mathbf{w}, b)$

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$



Linear  
Functions


# Linear Discriminant Function

- $g(\mathbf{x})$  is a linear function:

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$


- A hyper-plane in feature space  $\mathbf{x}$
- The unit-length normal vector of the hyper-plane:

$$\mathbf{n} = \frac{\mathbf{w}}{\|\mathbf{w}\|}$$




# Linear Discriminant Function

How can we classify the data using a linear discriminant in order to minimize the error rate ?




# Linear Discriminant Function

How can we classify the data using a linear discriminant in order to minimize the error rate ?



# Linear Discriminant Function

How can we classify the data using a linear discriminant in order to minimize the error rate ?



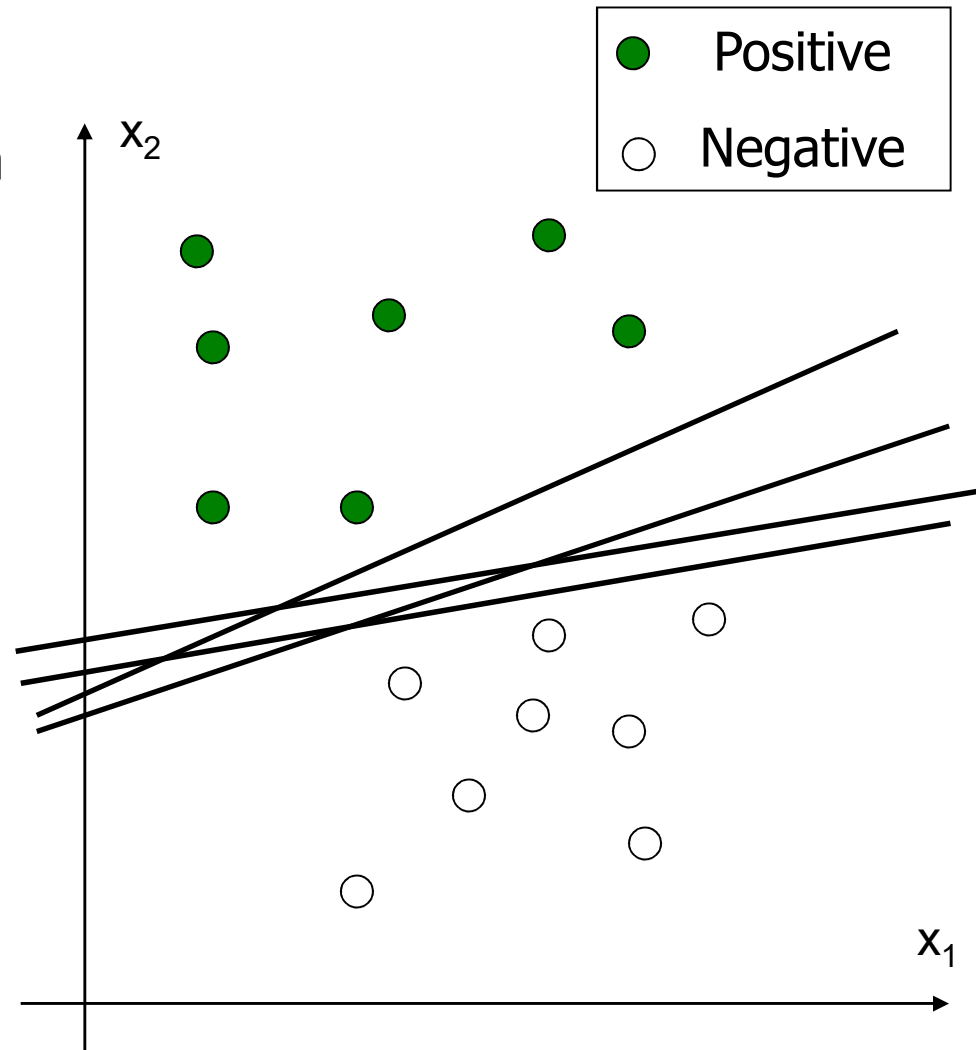


# Linear Discriminant Function

How can we classify the data using a linear discriminant in order to minimize the error rate ?

- Many possible answers!
- Which one is the best?

***A: The best is the one that gives the lowest error on new test data!***



# Large Margin Linear Classifier

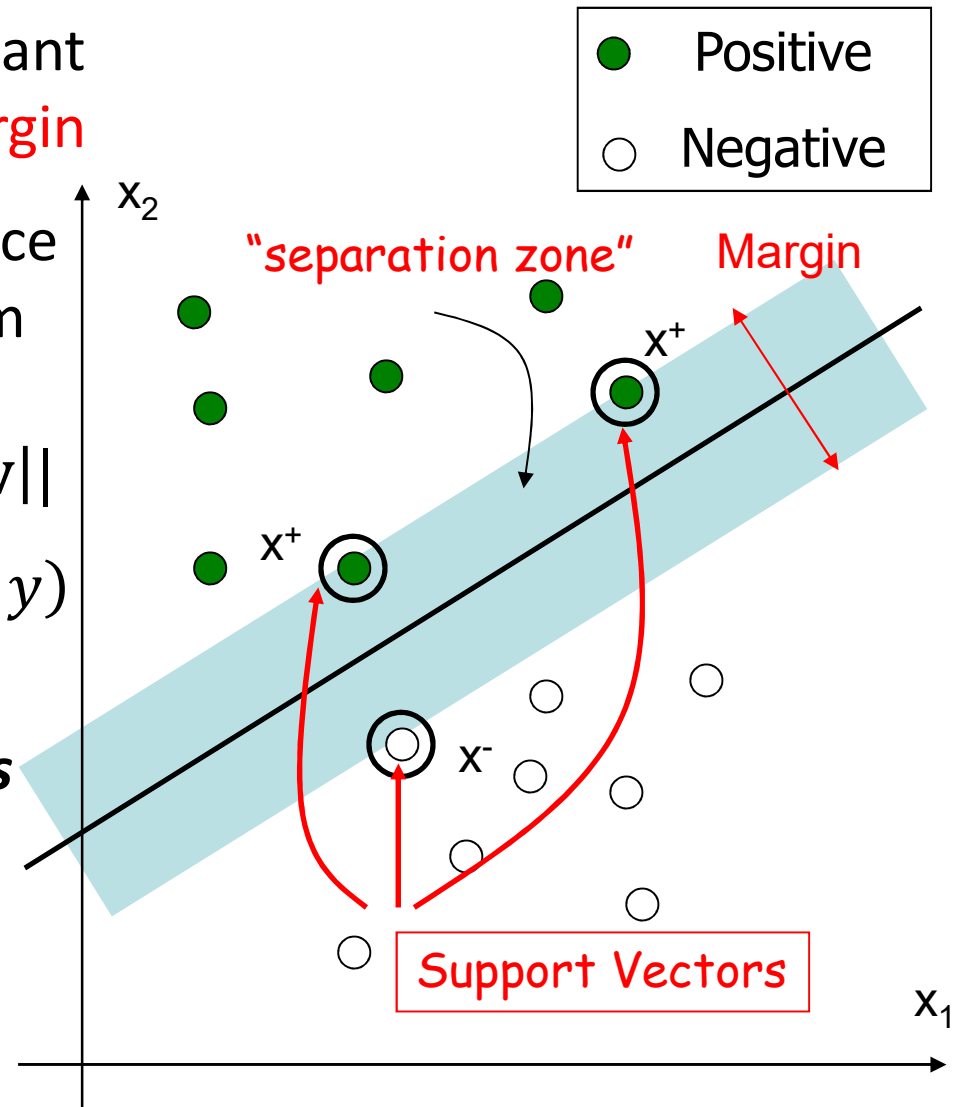
**One option:** the linear discriminant function with the maximum **margin**

- **Geometric margin** is the distance to a separating hyperplane from the point closest to it:

$$\rho_{\mathbf{w},b}(\mathbf{x}, y) = y(\mathbf{w}^T \mathbf{x} + b) / \|\mathbf{w}\|$$

$$\text{Margin } \rho_{\mathbf{w},b} := \min_{i=1,\dots,N} \rho_{\mathbf{w},b}(\mathbf{x}_i, y)$$

- Examples closest to the hyperplane are **support vectors**
- The **discriminant margin** is the maximum width of the band that can be drawn, separating contrastive support vectors



# Large Margin Linear Classifier

- Why is it good to focus on large margin?
  - Robust to outliers and thus strong generalization ability
- Fundamental result
  - If a d-dimensional dataset is enclosed inside a sphere of radius R, and the margin of the linear classifier is M, then the generalization error (difference between expected and empirical error) of the classifier is bounded by a function of the VC dimension, and

$$VC \leq \min \left\{ d, \left\lceil \frac{4R^2}{M^2} \right\rceil \right\} + 1$$


- In particular, independently from the dimension d, we can reduce the VC dimension by increasing the margin.

# VC Dimension

- Consider a binary classification problem, and a function class
- Each function of the class induces a labeling of patterns
- There are at most  $2^N$  labelings for  $N$  patterns
- If a very rich function class might be able to realize all  $2^N$  separations, it is said to **shatter** the  $N$  points
- However the function may not be rich enough
- The VC dimension is defined as the largest  $m$  such that there exist a set of  $m$  points which the class can shatter, or  $\infty$  if no such  $m$  exists
- It is a one number summary for the capacity of the learning machine

# VC Dimension Example

There exist VC dimension points that can be shattered, i.e. arbitrarily classified. Here: There exist 3 points in 2 dimensions that can be shattered.



**Figure 1.4** A simple VC dimension example. There are  $2^3 = 8$  ways of assigning 3 points to two classes. For the displayed points in  $\mathbb{R}^2$ , all 8 possibilities can be realized using separating hyperplanes, in other words, the function class can shatter 3 points. This would not work if we were given 4 points, no matter how we placed them. Therefore, the VC dimension of the class of separating hyperplanes in  $\mathbb{R}^2$  is 3.

# Cover's Theorem

- Gives the number of possible linear separations of  $N$  points, in general position, in a  $d$ -dimensional space
- If  $N \leq d + 1$  then  $2^N$  separations are possible  $VC \text{ dim} = d + 1$
- If  $N > d + 1$ , the number of linear separations is

$$2 \sum_{i=0}^d \binom{N-1}{i}$$

- As we increase  $d$ , there are more terms in the sum, VC ↗
- Points assumed in general position: however in practical applications points could be on lower-dimensional manifold

# Large Margin Linear Classifier

- Given a set of data points:  
 $\{(\mathbf{x}_i, y_i)\}, i = 1, 2, \dots, n$ , where

For  $y_i = +1$ ,  $\mathbf{w}^T \mathbf{x}_i + b > 0$


For  $y_i = -1$ ,  $\mathbf{w}^T \mathbf{x}_i + b < 0$

- Canonical Hyperplane**

Under a scale transformation on both  $\mathbf{w}$  and  $b$ , we can remove gauge in the above

For  $y_i = +1$ ,  $\mathbf{w}^T \mathbf{x}_i + b \geq 1$

For  $y_i = -1$ ,  $\mathbf{w}^T \mathbf{x}_i + b \leq -1$



# Large Margin Linear Classifier

- We know that


$$w^T x^+ + b = 1$$

$$w^T x^- + b = -1$$

- The separation is

$$S = n^T \cdot (x^+ - x^-) =$$

$$= \frac{w^T}{\|w\|} \cdot (x^+ - x^-) = \frac{2}{\|w\|}$$





# Large Margin Linear Classifier


- Formulation:

$$\max_{\mathbf{w}, b} \frac{2}{\|\mathbf{w}\|}$$

such that

$$\text{For } y_i = +1, \quad \mathbf{w}^T \mathbf{x}_i + b \geq 1$$

$$\text{For } y_i = -1, \quad \mathbf{w}^T \mathbf{x}_i + b \leq -1$$



# Large Margin Linear Classifier


- Formulation:

$$\min_{\mathbf{w}, b} \quad \frac{1}{2} \|\mathbf{w}\|^2$$

such that

$$\text{For } y_i = +1, \quad \mathbf{w}^T \mathbf{x}_i + b \geq 1$$

$$\text{For } y_i = -1, \quad \mathbf{w}^T \mathbf{x}_i + b \leq -1$$



# Large Margin Linear Classifier


- Formulation:

$$\min_{\mathbf{w}, b} \quad \frac{1}{2} \|\mathbf{w}\|^2$$

such that

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

- Quadratic program with linear constraints




# Solving the Optimization Problem

Quadratic  
programming  
with linear  
constraints

$$\begin{aligned} \min_{\mathbf{w}, b} \quad & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{s.t.} \quad & y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \end{aligned}$$

Lagrangian  
Function



$$\begin{aligned} \min_{\mathbf{w}, b} \max_{\alpha} \quad & L_p(\mathbf{w}, b, \alpha_i) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i (y_i (\mathbf{w}^T \mathbf{x}_i + b) - 1) \\ \text{s.t.} \quad & \alpha_i \geq 0 \end{aligned}$$

The Lagrangian needs to be minimized w.r.t.  $\mathbf{w}, b$ , and maximized w.r.t  $\alpha_i$

# Solving the Optimization Problem

$$\begin{aligned} \min_{\mathbf{w}, b} \max_{\alpha} L_p(\mathbf{w}, b, \alpha_i) &= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i (y_i (\mathbf{w}^T \mathbf{x}_i + b) - 1) \\ \text{s.t.} \quad \alpha_i &\geq 0 \end{aligned}$$

$$\frac{\partial L_p}{\partial \mathbf{w}} = 0$$



$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

Solution is an expansion in terms of training examples

$$\frac{\partial L_p}{\partial b} = 0$$




$$\sum_{i=1}^n \alpha_i y_i = 0$$

Due to strict convexity,  $\mathbf{w}$  is unique although  $\alpha_i$ 's need not be

# Solving the Optimization Problem

$$\begin{aligned} \min_{\mathbf{w}, b} \max_{\alpha} L_p(\mathbf{w}, b, \alpha_i) &= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i (y_i (\mathbf{w}^T \mathbf{x}_i + b) - 1) \\ \text{s.t.} \quad &\alpha_i \geq 0 \end{aligned}$$

Lagrangian Dual  
Problem



$$\begin{aligned} \text{maximize} \quad & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \\ \max_{\alpha} \quad & \\ \text{s.t.} \quad & \alpha_i \geq 0, \text{ and } \sum_{i=1}^n \alpha_i y_i = 0 \end{aligned}$$

Convex quadratic optimization problem. Using a QP solver, gives us the uniquely best  $\alpha_i$ .

# Solving the Optimization Problem


- Suppose we have found the optimal  $\alpha$
- From the KKT conditions, we know:

$$\alpha_i (y_i (\mathbf{w}^T \mathbf{x}_i + b) - 1) = 0$$

- Thus, only support vectors have  $\alpha_i \neq 0$
- The solution has the form:

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i = \sum_{i \in \text{SV}} \alpha_i y_i \mathbf{x}_i$$

We can get  $b$  from  $y_i (\mathbf{w}^T \mathbf{x}_i + b) - 1 = 0$   
where  $\mathbf{x}_i$  is any support vector



# Solving for $b$

For support vectors

$$y_s(\mathbf{w}^\top \cdot \mathbf{x}_s + b) = 1$$

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

$$y_s \left( \sum_{i \in S} \alpha_i y_i \mathbf{x}_i^\top \cdot \mathbf{x}_s + b \right) = 1 \quad | \quad \mathbf{x} \quad y_s$$

$$b = \frac{1}{N_{SV}} \sum_{i=1}^{N_{SV}} (y_s - \sum_{i \in S} \alpha_i y_i \mathbf{x}_i^\top \cdot \mathbf{x}_s)$$

*Better take an average over support vectors*



# Solving the Optimization Problem

- The linear discriminant function is:

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = \sum_{i \in SV} \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$


- Relies on a *dot product* between the test point  $\mathbf{x}$  and the support vectors  $\mathbf{x}_i$
- Solving the optimization problem involved computing the **dot products**

between all pairs of training points

- Negative side: with many points, this is expensive
- Positive side: The algorithm and solution only needs this matrix of products from the training points, not the points itself. We will take advantage of that.

# `Soft Margin' Linear Classifier

- What if data is not linear separable due to noise or outliers?
- Slack variables  $\xi_i$  can be added to allow for the mis-classification of difficult or noisy data



# `Soft Margin' Linear Classifier

- Formulation:

$$\min_{\mathbf{w}, b, \xi} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

such that

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i$$

$$\xi_i \geq 0$$

- for  $0 \leq \xi \leq 1$ , point is between margin and correct side of hyperplane
- for  $\xi > 1$ , point is misclassified

- Parameter  $C$  can be viewed as a means to control over-fitting
  - small  $C$  allows constraints to be easily ignored: *large margin*
  - large  $C$  makes constraints hard to ignore: *narrow margin*
  - $C = \infty$  enforces all constraints: *hard margin*

# `Soft Margin' Linear Classifier

- Formulation (Lagrangian Dual Problem)

$$\max_{\boldsymbol{\alpha}} \quad \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

such that

$$0 \leq \alpha_i \leq C$$

$$\sum_{i=1}^n \alpha_i y_i = 0$$

# `Soft Margin' Interpretation (I)

- The constraint  $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i$  can be written more concisely as

$$y_i g(\mathbf{x}_i) \geq 1 - \xi_i \Leftrightarrow \xi_i = \max(0, 1 - y_i g(\mathbf{x}_i))$$

- Hence we need to solve the learning problem

$$\min_{\mathbf{w}, b} ||\mathbf{w}'||^2 + C \sum_{i=1}^N \max(0, 1 - y_i g(\mathbf{x}_i))$$


# `Soft Margin' Interpretation (II)

We need to solve the learning problem

$$\min_{\mathbf{w}, b} ||\mathbf{w}'||^2 + C \sum_{i=1}^N \max(0, 1 - y_i g(\mathbf{x}_i))$$

- $y_i g(\mathbf{x}_i) > 1 \Rightarrow$  point is outside margin and does not contribute to loss
- $y_i g(\mathbf{x}_i) = 1 \Rightarrow$  point is on margin and does not contribute to loss (as in hard margin)
- $y_i g(\mathbf{x}_i) < 1 \Rightarrow$  point violates margin constraint and contributes to loss


# SVM uses Hinge Loss




Can be viewed as an approximation to the 0-1 loss

# Non-linear SVMs


- Datasets that are linearly separable with noise work out great:



- But what are we going to do if the dataset is just too hard?




- How about... mapping data to a higher-dimensional space:





# Non-linear SVMs: Feature Space

General idea: the original input space can be mapped to some higher-dimensional feature space where the training set is separable



# How to Use the Feature Space?

- The feature point  $\mathbf{z} = \phi(\mathbf{x})$  corresponding to an input point  $\mathbf{x}$  is called the image (or the lifting) of  $\mathbf{x}$  ; the input point  $\mathbf{x}$  , if any, corresponding to a given feature vector  $\mathbf{z}$  is called the pre-image of  $\mathbf{z}$
- The naive way to use a feature space is to explicitly compute the image of every training and testing point, and run algorithm fully in feature space
- Two potential problems
  - The feature space may be very high dimensional or infinite dimensional, so direct (explicit) calculations in such feature space may not be practical, or even possible
  - We may sometimes want to map back an answer from feature space to the input space. This is called the pre-image problem. For some kernels, analytical expressions are available, but in most other cases some form of (local) optimization may be necessary

# Nonlinear SVMs: The Kernel Trick

- With this mapping, our discriminant function is now:

$$g(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b = \sum_{i \in \text{SV}} \alpha_i \phi(\mathbf{x}_i)^T \phi(\mathbf{x}) + b$$

# Nonlinear SVMs: The Kernel Trick

- With this mapping, our discriminant function is now:

$$g(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b = \sum_{i \in \text{SV}} \alpha_i \boxed{\phi(\mathbf{x}_i)^T \phi(\mathbf{x})} + b$$

- No need to know this mapping explicitly, because we only use the *dot product* of feature vectors both in training and in testing

# Nonlinear SVMs: The Kernel Trick

- With this mapping, our discriminant function is now:

$$g(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b = \sum_{i \in \text{SV}} \alpha_i \boxed{\phi(\mathbf{x}_i)^T \phi(\mathbf{x})} + b$$

- No need to know this mapping explicitly, because we only use the *dot product* of feature vectors both in training and in testing
- A *kernel function* is defined as a function that corresponds to a dot product of two feature vectors in some expanded feature space:

$$K(\mathbf{x}_i, \mathbf{x}_j) \equiv \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

- Instead of specifying and computing features, can define and compute kernel only.

# Positive Definite Kernels

- **Gram Matrix.** Given a function  $k: X^2 \rightarrow \mathbf{R}$  (or  $\mathbf{C}$ ), and patterns  $x_1, \dots, x_m \in X$ , the  $m \times m$  matrix  $K$  with elements  $K_{ij} := k(x_i, x_j)$  is called the Gram matrix (or kernel matrix) of  $k$  w.r.t  $x_1, \dots, x_m$ .
- **Positive definite kernel.** A complex  $m \times m$  matrix  $K$  satisfying  $\sum_{ij} c_i \bar{c}_j K_{ij} \geq 0, \forall c_i \in \mathbf{C}$  is called positive definite. Similarly, a real symmetric  $m \times m$  matrix  $K$  satisfying the above for all  $c_i \in \mathbf{R}$  is called positive definite.

positive definite kernels  $\equiv$  Mercer kernels  $\equiv$  reproducing kernels  $\equiv$  admissible kernels  $\equiv$  support vector kernels  $\equiv$  covariance functions

# Examples of Kernels

Examples of commonly-used kernel functions:

- Linear kernel:  $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$
- Polynomial kernel:  $K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^p$
- Gaussian (Radial-Basis Function (RBF) ) kernel:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$

- Sigmoid:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \tanh(\beta_0 \mathbf{x}_i^T \mathbf{x}_j + \beta_1)$$

# Generality of Kernel Trick

- Given an algorithm expressed in terms of a positive-definite kernel  $k$ , we can construct an alternative algorithm by replacing  $k$  with another positive-definite kernel  $\tilde{k}$
- This is not limited to only cases when  $k$  is a dot product in the input domain
- Any algorithm that only depends on dot products (i.e. is rotationally invariant) can be kernelized
- Kernels are defined on general sets (rather than just dot product spaces!) and their use leads to an embedding of general data types in linear spaces



# Nonlinear SVM: Optimization

- Formulation (Lagrangian Dual Problem)

$$\max_{\alpha} \quad \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j)$$

such that

$$0 \leq \alpha_i \leq C$$

$$\sum_{i=1}^n \alpha_i y_i = 0$$

- The solution of the discriminant function is

$$g(\mathbf{x}) = \sum_{i \in SV} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b$$

# Support Vector Machine: Algorithm

1. Choose a kernel function
2. Choose a value for  $C$
3. Solve the quadratic programming problem  
(many software packages available, e.g. libsvm)
4. Construct the discriminant function from the support vectors

# Support Vector Machines

- A support vector machine (SVM) is nothing more than a kernelized maximum-margin hyperplane classifier
- You train it by solving the dual quadratic programming problem
- You run it by evaluating the kernel function between the test point and each of the “active” training points, called support vectors
- This combination of (1) kernel trick, (2) maximum margin (minimum norm) and (3) the resulting sparsity has turned out to be very effective and popular
- In practice, the hard part from a learning point of view is selecting the kernel function (there is a lot of research on this) and from a computational point of view it is solving the large QP efficiently (most popular method: Sequential Minimal Optimization, estimates pairs of parameters sequentially)

# SVM Applet Demo

<https://cs.stanford.edu/people/karpathy/svmjs/demo/>

# Properties of Kernels

- Kernels are symmetric in their arguments:

$$K(\mathbf{x}_1, \mathbf{x}_2) = K(\mathbf{x}_2, \mathbf{x}_1)$$

- They are positive valued for any inputs:  $K(\mathbf{x}_1, \mathbf{x}_2) \geq 0$

- The Cauchy-Schwartz inequality holds:

$$K^2(\mathbf{x}_1, \mathbf{x}_2) \leq K(\mathbf{x}_1, \mathbf{x}_1)K(\mathbf{x}_2, \mathbf{x}_2)$$

- Technically, to use a function as a kernel, it must satisfy Mercer's conditions for a positive-definite operator

- The intuition is easy to grasp for finite spaces

- Discretize  $\mathbf{x}$  space as densely as desired into buckets  $\mathbf{x}_i$
- Between each two cells  $\mathbf{x}_i, \mathbf{x}_j$ , compute the kernel function, and write these values as a (symmetric) matrix  $M_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$
- If the matrix is positive definite, the kernel is OK

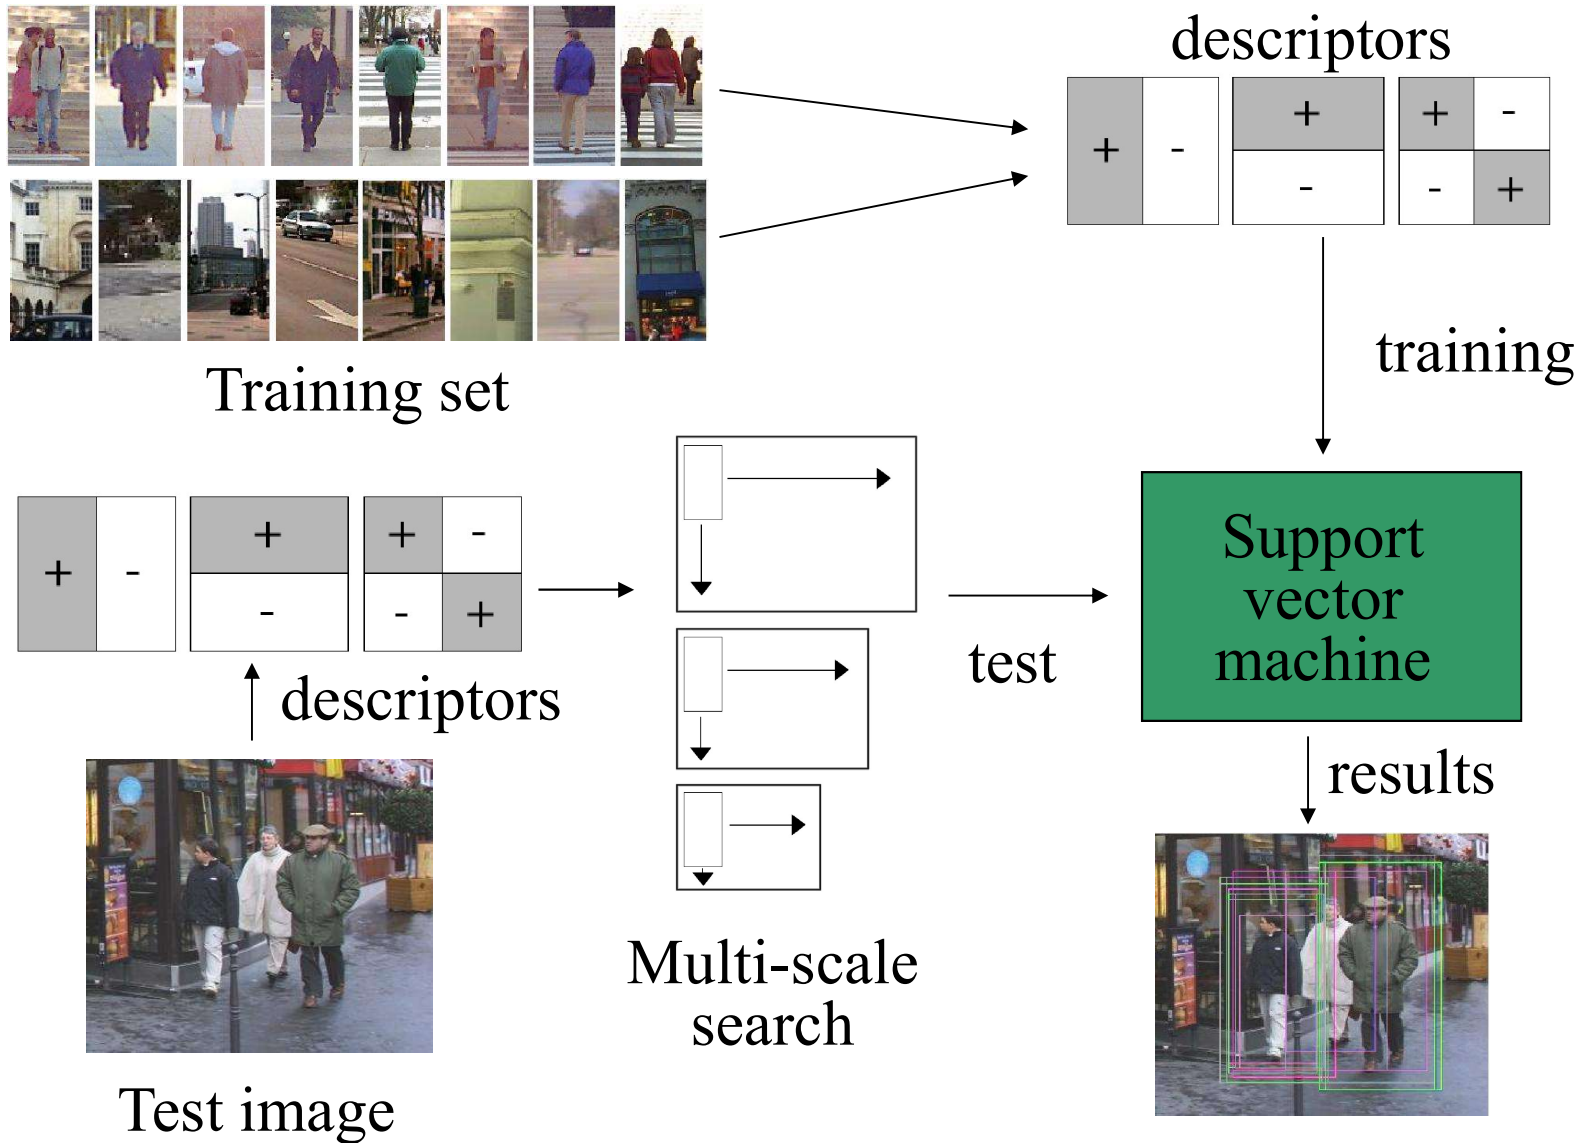
# Kernel Closure Rules

**Very useful for designing new kernels from existing kernels**


- The sum of any two kernels is a kernel
- The product of any two kernels is a kernel
- A kernel plus a constant is a kernel
- A scalar times a kernel is a kernel

# Support Vector Machine Detector

*Papageorgiou, Poggio*



# Video: Pedestrian Detection





# Scalability Issues

Although we circumvented infinite dimensionality,

$$f(\mathbf{x}) = \sum_i \alpha_i K(\mathbf{x}, \mathbf{x}_i)$$

- In training:
  - # optimization variables = # training examples  $N$
- In testing:
  - Need to evaluate kernel between test data and each training example
- Training and testing for millions of examples unfeasible
  - e.g. in ImageNet need to classify 9 million images

# Linear versus Kernel Methods

	Linear	Kernel
Model	$f(x) = w^T x$	$f(x) = \sum_i \alpha_i k(x, x_i)$
Number of optimization variables	Input dimensionality $d$	# training examples $N$
Training time	$O(Nd^2)$	$O(N^2d) \sim O(N^3d)$
Testing time	$O(d)$	$O(Nd)$
<a href="#">Caltech-101</a> Accuracy (BOW feature)	49% (Vedaldi and Zisserman 2010)	64% (Vedaldi and Zisserman 2010)
<a href="#">Caltech-101</a> Accuracy (multiple kernels)	N/A	82% (Gehler & Nowozin 2009) (Li et al. 2010)

*Good things are worth doing slowly*